


# Lecture 10-11

## Artificial Intelligence

Khola Naseem  
khola.naseem@uet.edu.pk




# Adversarial Search



- Multi agent environments : any given agent will need to consider the actions of other agents and how they affect its own welfare.
- The unpredictability of these other agents can introduce many possible contingencies

# Adversarial Search

- Examine the problems that arise when we try to plan ahead in a world where other agents are planning against us.
- Competitive environments, in which the agent's goals are in conflict require **adversarial search**
- A good example is in board games, Chess, Tic-tac-toe, backgammon and many more



# Typical AI assumptions

- AI games are a specialized kind - deterministic, turn taking, two-player, zero sum games of perfect information
- a **zero-sum game** is a mathematical representation of a situation in which a participant's gain (or loss) of utility is exactly balanced by the losses (or gains) of the utility of other participant(s)
- In our terminology – deterministic, fully observable environments with two agents whose actions alternate and the utility values at the end of the game are always equal and opposite (+1 and -1)
- If a player wins a game of chess (+1), the other player necessarily loses (-1)

# Search vs games

## ➤ Search – no adversary

- Solution is (heuristic) method for finding goal
- Heuristic techniques can find optimal solution
- Evaluation function: estimate of cost from start to goal through given node
- Examples: path planning, scheduling activities

## ➤ Games – adversary

- Solution is strategy (strategy specifies move for every possible opponent reply).
- Optimality depends on opponent. Why?
- Time limits force an approximate solution
- Evaluation function: evaluate “goodness” of game position
- Examples: chess, checkers

# Size of search tree

- Games, unlike most of the toy problems are interesting because they are too hard to solve.
  - $b$  = branching factor
  - $d$  = number of moves by both players
  - Search tree is  $O(bd)$
- For example, chess has an average branching factor of about 35, and games often go to 50 moves by each player, so the search tree has about  $35^{100}$  nodes

# Games step up

- Two players: MAX and MIN
- MAX moves first and they take turns until the game is over
  - Winner gets award, loser gets penalty.
- Games as search:
  - Initial state( $S_0$ ): The initial state, which specifies how the game is set up at the start.
  - PLAYER(s): Defines which player has the move in a state.
  - ACTIONS(s): Returns the set of legal moves in a state.
  - RESULT(s, a): The transition model, which defines the result of a move.
  - TERMINAL-TEST(s): A terminal test, which is true when the game is over and false otherwise. States where the game has ended are called terminal states.
  - UTILITY(s, p): A utility function (also called an objective function or payoff function), defines the final numeric value for a game that ends in terminal state s for a player p. In chess, the outcome is a win, loss, or draw, with values +1, 0, or 1/2 .


# Deterministic Games

- Many possible formalizations, one is:
  - States:  $S$  (start at  $s_0$ )
  - Players:  $P = \{1 \dots N\}$  (usually take turns)
  - Actions:  $A$  (may depend on player / state)
  - Transition Function:  $S \times A \rightarrow S$
  - Terminal Test:  $S \rightarrow \{t, f\}$
  - Terminal Utilities:  $S \times P \rightarrow R$




# Partial Game Tree for Tic-Tac-Toe

- The initial state, ACTIONS function, and RESULT function define the game tree for the game—a tree where the nodes are game states and the edges are moves



# Partial Game Tree for Tic-Tac-Toe

- The initial state, ACTIONS function, and RESULT function define the game tree for the game—a tree where the nodes are game states and the edges are moves




- For tic-tac-toe the game tree is relatively small—fewer than  $9! = 362,880$  terminal nodes.

# Minimax Algorithm:

- Find the optimal strategy for MAX assuming an infallible MIN opponent
  - Need to compute this all the down the tree
- Assumption: Both players play optimally!
- Given a game tree, the optimal strategy can be determined by using the minimax value of each node.


# Minimax Algorithm:

- Two player game tree:




# Minimax Algorithm:

- Two player game tree:



# Minimax Algorithm:

- Two player game tree:



# Minimax Algorithm:

- **What if MIN does not play optimally?**
- Definition of optimal play for MAX assumes MIN plays optimally:
  - maximizes worst-case outcome for MAX
- But if MIN does not play optimally, MAX will do even better
  - Can prove this

# Minimax Algorithm:

## ➤ Goal

MINIMAX( $s$ ) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$



# Minimax Algorithm:


## ➤ Pseudocode for Minimax Algorithm

```
function minimax(position, depth, maximizingPlayer)
    if depth == 0 or game over in position
        return static evaluation of position

    if maximizingPlayer
        maxEval = -infinity
        for each child of position
            eval = minimax(child, depth - 1, false)
            maxEval = max(maxEval, eval)
        return maxEval


    else
        minEval = +infinity
        for each child of position
            eval = minimax(child, depth - 1, true)
            minEval = min(minEval, eval)
        return minEval

// initial call
minimax(currentPosition, 3, true)
```




# Minimax Algorithm:

- Two player game tree:



# Minimax Algorithm:

➤ Two player game tree:



```
function minimax(position, depth, maximizingPlayer)
    if depth == 0 or game over in position
        return static evaluation of position

    if maximizingPlayer
        maxEval = -infinity
        for each child of position
            eval = minimax(child, depth - 1, false)
            maxEval = max(maxEval, eval)
        return maxEval

    else
        minEval = +infinity
        for each child of position
            eval = minimax(child, depth - 1, true)
            minEval = min(minEval, eval)
        return minEval

// initial call
minimax(currentPosition, 3, true)
```

# Minimax Algorithm:

- Complete?

- Yes (if tree is finite).

- Optimal?

- Yes (against an optimal opponent).

- Can it be beaten by an opponent playing sub-optimally?

- Time complexity?

- $O(b^m)$

- Space complexity?


- $O(bm)$  (depth-first search, generate all actions at once)

# Minimax Algorithm:

- Complete depth-first exploration of the game tree
- Assumptions:
  - Max depth =  $m$ ,  $b$  legal moves at each point
- then the time complexity of the minimax algorithm is  $O(b^m)$ . The space complexity is  $O(bm)$  for an algorithm that generates all actions at once
- For real games, of course, the time cost is totally impractical, but this algorithm serves as the basis for the mathematical analysis of games and for more practical algorithms.

# Multiplayer games:

- First, we need to replace the single value for each node with a vector of values. For example, in a three-player game with players A, B, and C, a vector  $(v_A, v_B, v_C)$  is associated with each node.
- For terminal states, this vector gives the utility of the state from each player's viewpoint.



# Multiplayer games:

- **Zero sum games:** zero-sum describes a situation in which a participant's gain or loss is exactly balanced by the losses or gains of the other participant(s).
- If the total gains of the participants are added up, and the total losses are subtracted, they will sum to zero

	B chooses B1	B chooses B2	B chooses B3
A chooses A1	+3	-2	+2
A chooses A2	-1	0	+4
A chooses A3	-4	-3	+1

# Multiplayer games:

## ➤ **Alliance:**

- suppose A and B are in weak positions and C is in a stronger position.
- In this way, collaboration emerges from purely selfish behavior.
- as soon as C weakens the alliance loses its value, and either A or B could violate the agreement. In some cases, explicit alliances merely make concrete



# Minimax Issue:


- Number of game states is exponential in the number of moves.
- Solution: Do not examine every node
  - => pruning
  - Remove branches that do not influence final decision

# (Static) Heuristic Evaluation Functions:

- An Evaluation Function:
  - Estimates how good the current board configuration is for a player.
  - Typically, evaluate how good it is for the player, how good it is for the opponent, then subtract the opponent's score from the player's.
  - Often called “static” because it is called on a static board position.
  - Chess: Value of all white pieces - Value of all black pieces
- Typical values from -infinity (loss) to +infinity (win) or  $[-1, +1]$ .
- If the board evaluation is  $X$  for a player, it's  $-X$  for the opponent
  - “Zero-sum game”

# (Static) Heuristic Evaluation Functions:

- material value for each piece: each pawn is worth 1, a knight or bishop is worth 3, a rook 5, and the queen 9 and so on
- Chess: Value of all white pieces - Value of all black pieces



For chess, typically *linear* weighted sum of *features*


$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g.,  $w_1 = 9$  with

$f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$


# (Static) Heuristic Evaluation Functions:

## Evaluation functions



Black to move

White slightly better



White to move

Black winning

For chess, typically *linear* weighted sum of *features*

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g.,  $w_1 = 9$  with

$f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$

# Alpha-beta pruning:


- The problem with minimax search is that the number of game states it has to examine is exponential in the depth of the tree.
- we can't eliminate the exponent, but it turns out we can effectively cut it in half.
- It is possible to compute the exact minimax decision without expanding every node in the game tree
- That is, we can borrow the idea of pruning the particular technique we examine is called alpha–beta pruning
- Prune away branches that cannot possibly influence the final decision.

# Alpha-beta pruning:

- Depth first search – only considers nodes along a single path at any time
- $\alpha$  = highest-value choice that we can guarantee for MAX so far in the current subtree.
- $\beta$  = lowest-value choice that we can guarantee for MIN so far in the current subtree.
- update values of  $\alpha$  and  $\beta$  during search and prunes remaining branches as soon as the value is known to be worse than the current  $\alpha$  or  $\beta$  value for MAX or MIN.


# Alpha-beta pruning:

- It is possible to compute the exact minimax decision without expanding every node in the game tree



# Alpha-beta pruning:

- It is possible to compute the exact minimax decision without expanding every node in the game tree






# Alpha-beta pruning:

- When to Prune
- Prune whenever  $\alpha \geq \beta$
- Prune below a Max node whose alpha value becomes greater than or equal to the beta value of its ancestors.  $\alpha \geq \beta$ 
  - Max nodes update alpha based on children's returned values.
- Prune below a Min node whose beta value becomes less than or equal to the alpha value of its ancestors.  $\beta \leq \alpha$ 
  - Min nodes update beta based on children's returned values.


# Alpha-beta pruning:

- It is possible to compute the exact minimax decision without expanding every node in the game tree




# Alpha-beta pruning:

- It is possible to compute the exact minimax decision without expanding every node in the game tree




# Alpha-beta pruning:

- It is possible to compute the exact minimax decision without expanding every node in the game tree




# Alpha-beta pruning:

- It is possible to compute the exact minimax decision without expanding every node in the game tree




# Alpha-beta pruning:

- It is possible to compute the exact minimax decision without expanding every node in the game tree




# Alpha-beta pruning:

- It is possible to compute the exact minimax decision without expanding every node in the game tree



# Alpha-beta pruning:


- It is possible to compute the exact minimax decision without expanding every node in the game tree





# Alpha-beta pruning:

- It is possible to compute the exact minimax decision without expanding every node in the game tree




# Alpha-beta pruning:

- In practice often get  $O(b^{(m/2)})$  rather than  $O(b^m)$
- Worst-Case
  - branches are ordered so that no pruning takes place. In this case alpha-beta gives no improvement over exhaustive search

# Alpha-beta pruning:

➤ Example:



# Alpha-beta pruning:

➤ Example:

**function** ALPHA-BETA-SEARCH(*state*) **returns** an action  
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
**return** the *action* in ACTIONS(*state*) with value *v*

---


**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
 $v \leftarrow -\infty$   
**for each** *a* **in** ACTIONS(*state*) **do**  
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
**if**  $v \geq \beta$  **then return** *v*  
 $\alpha \leftarrow \text{MAX}(\alpha, v)$   
**return** *v*

---


**function** MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
 $v \leftarrow +\infty$   
**for each** *a* **in** ACTIONS(*state*) **do**  
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
**if**  $v \leq \alpha$  **then return** *v*  
 $\beta \leftarrow \text{MIN}(\beta, v)$   
**return** *v*

# Alpha-beta pruning:


➤ Example:



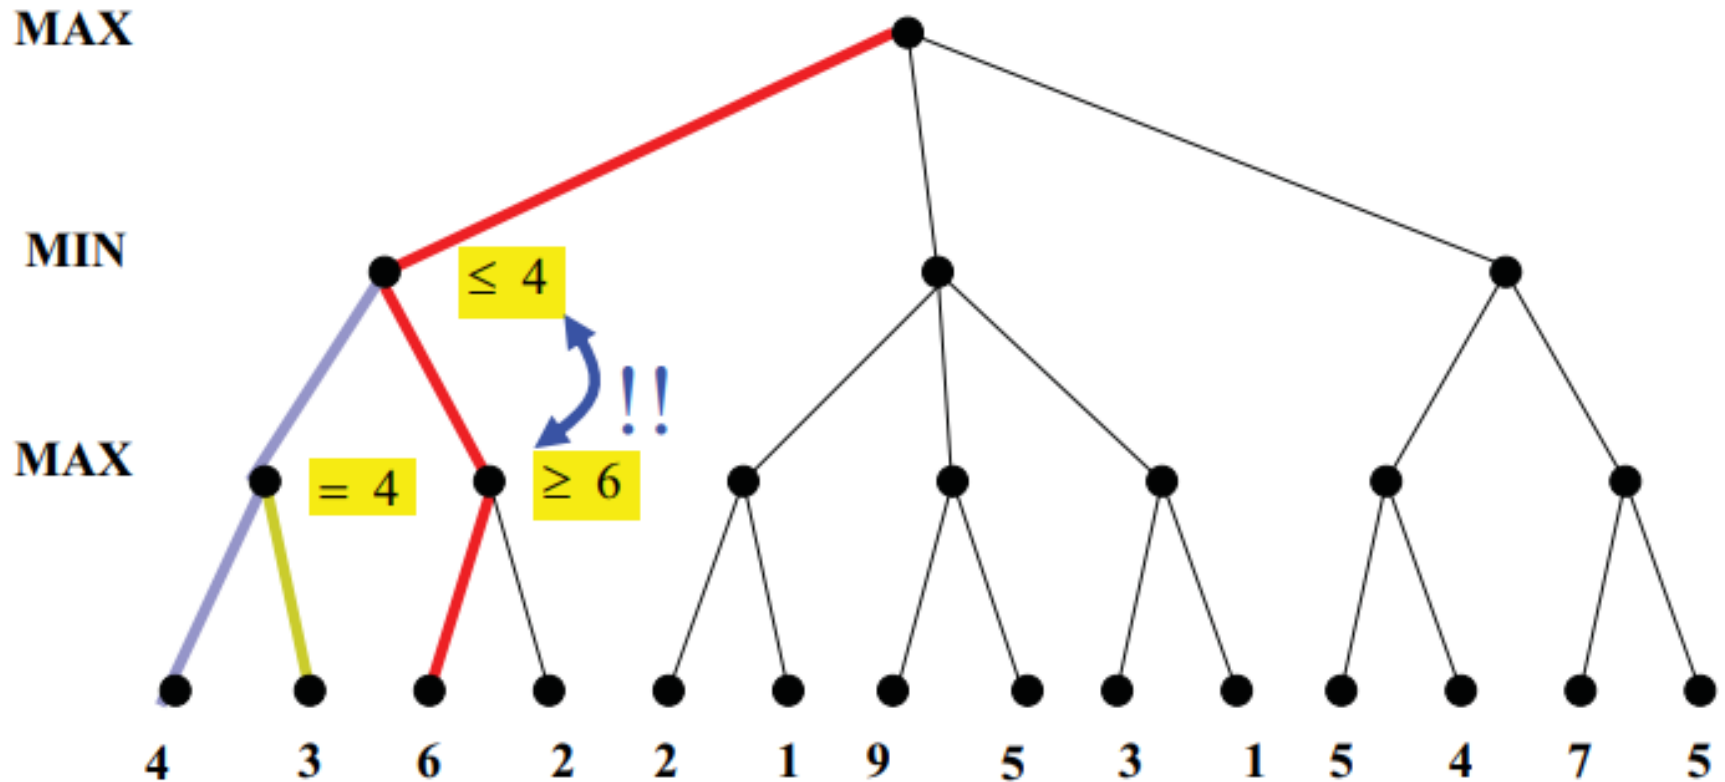
## Alpha beta pruning. Example



## Alpha beta pruning. Example




## Alpha beta pruning. Example

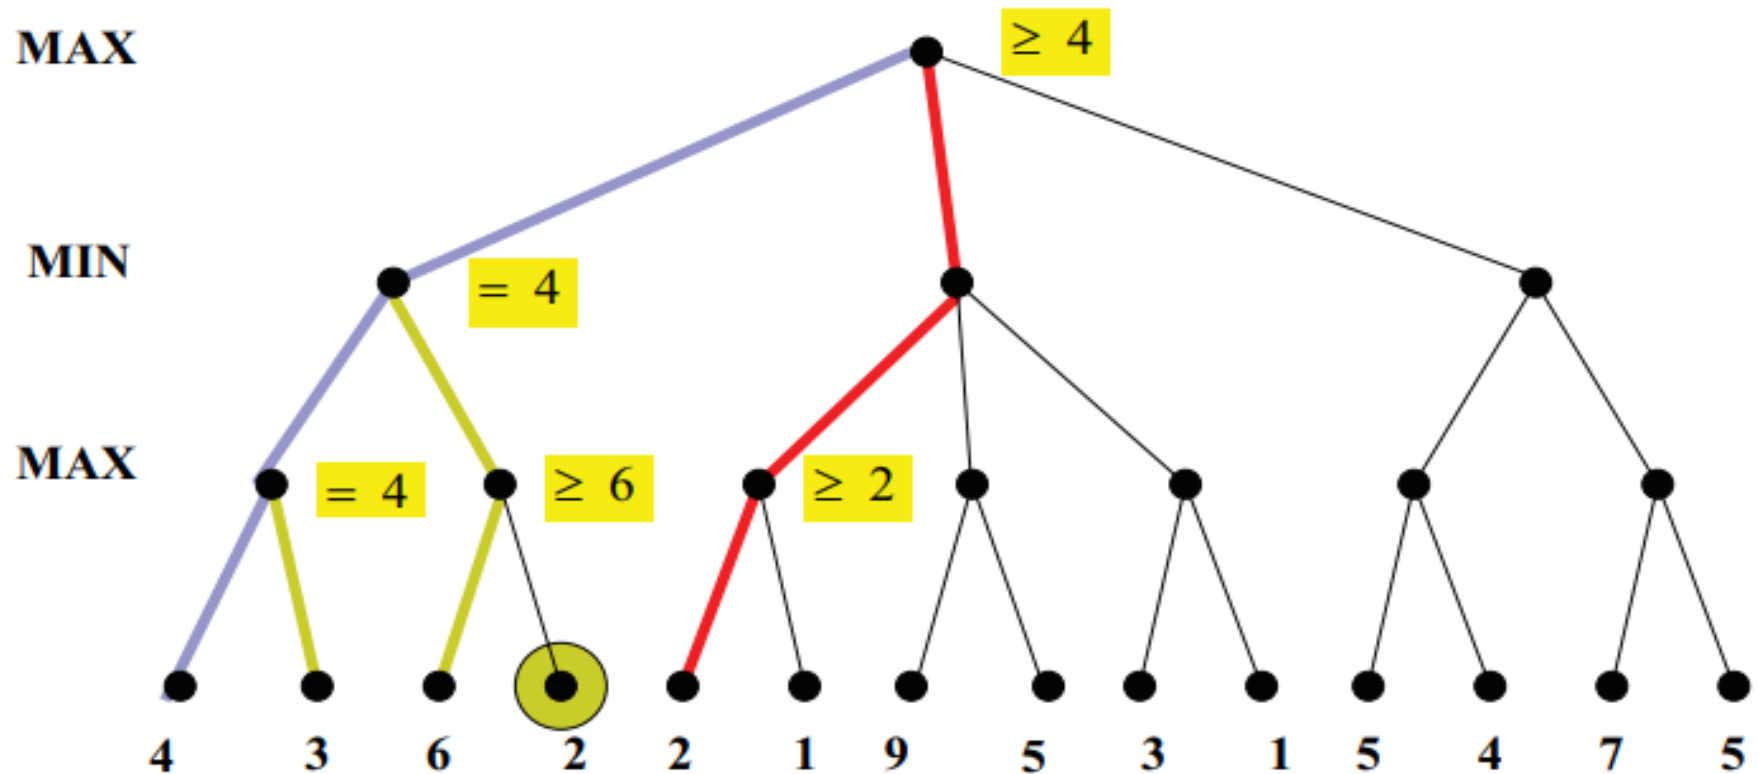




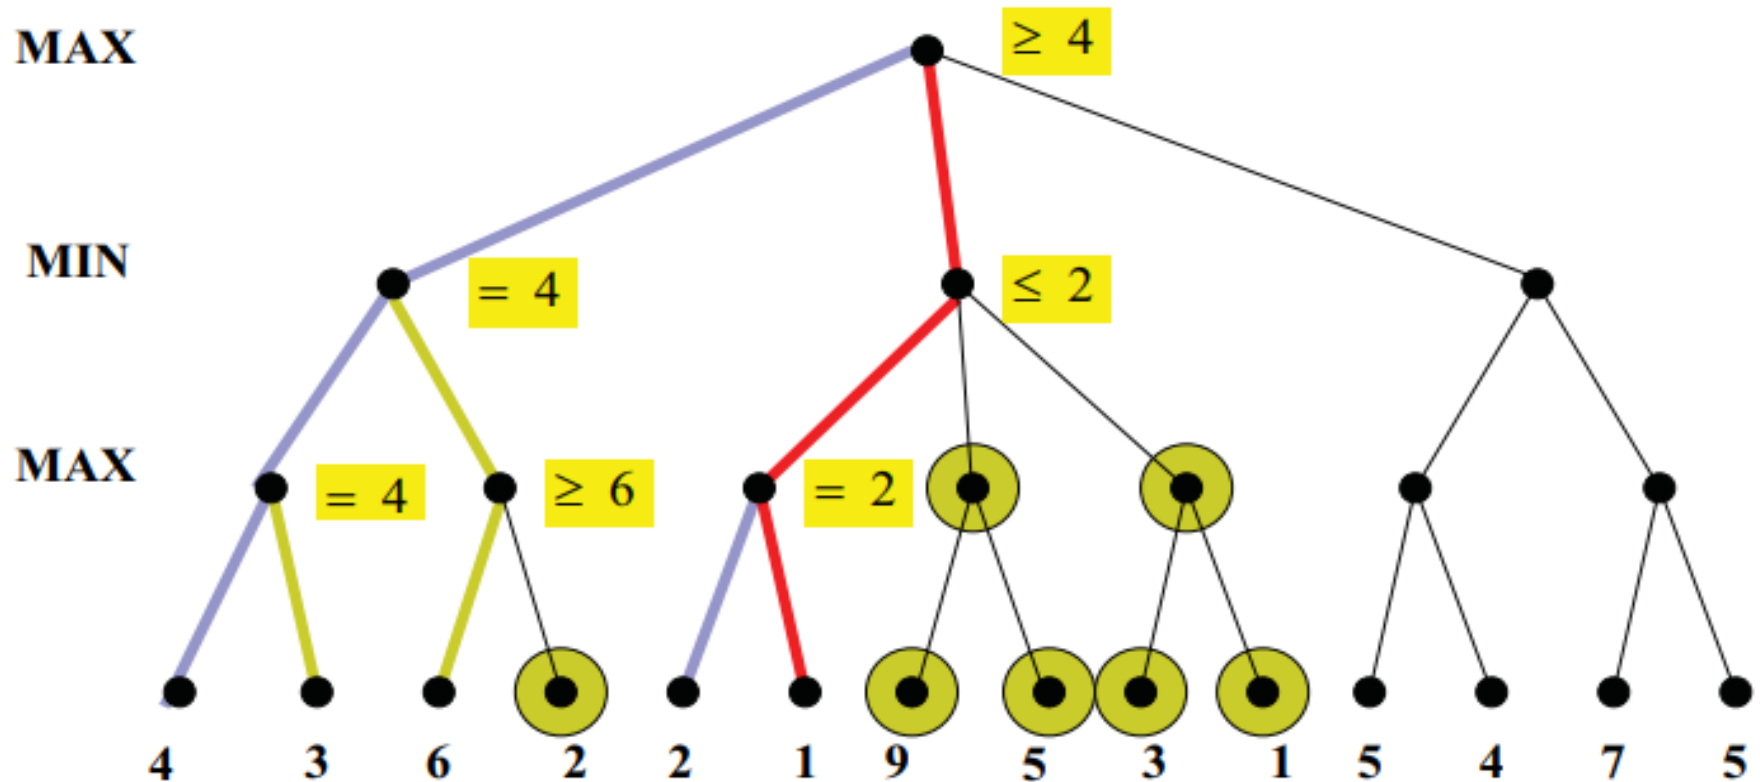
# Alpha beta pruning. Example



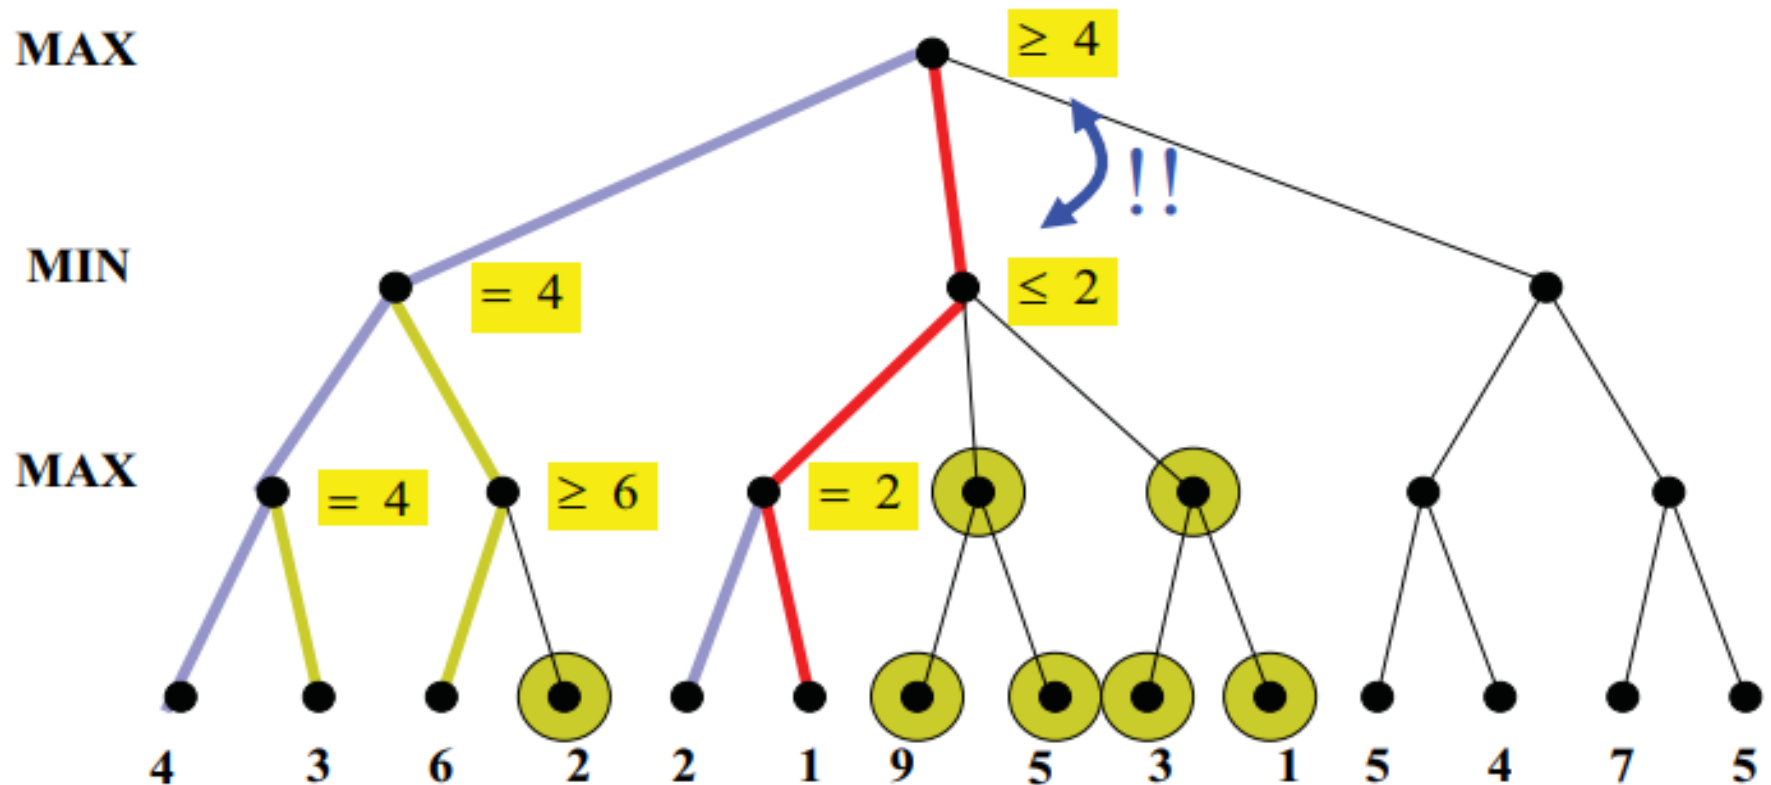
## Alpha beta pruning. Example



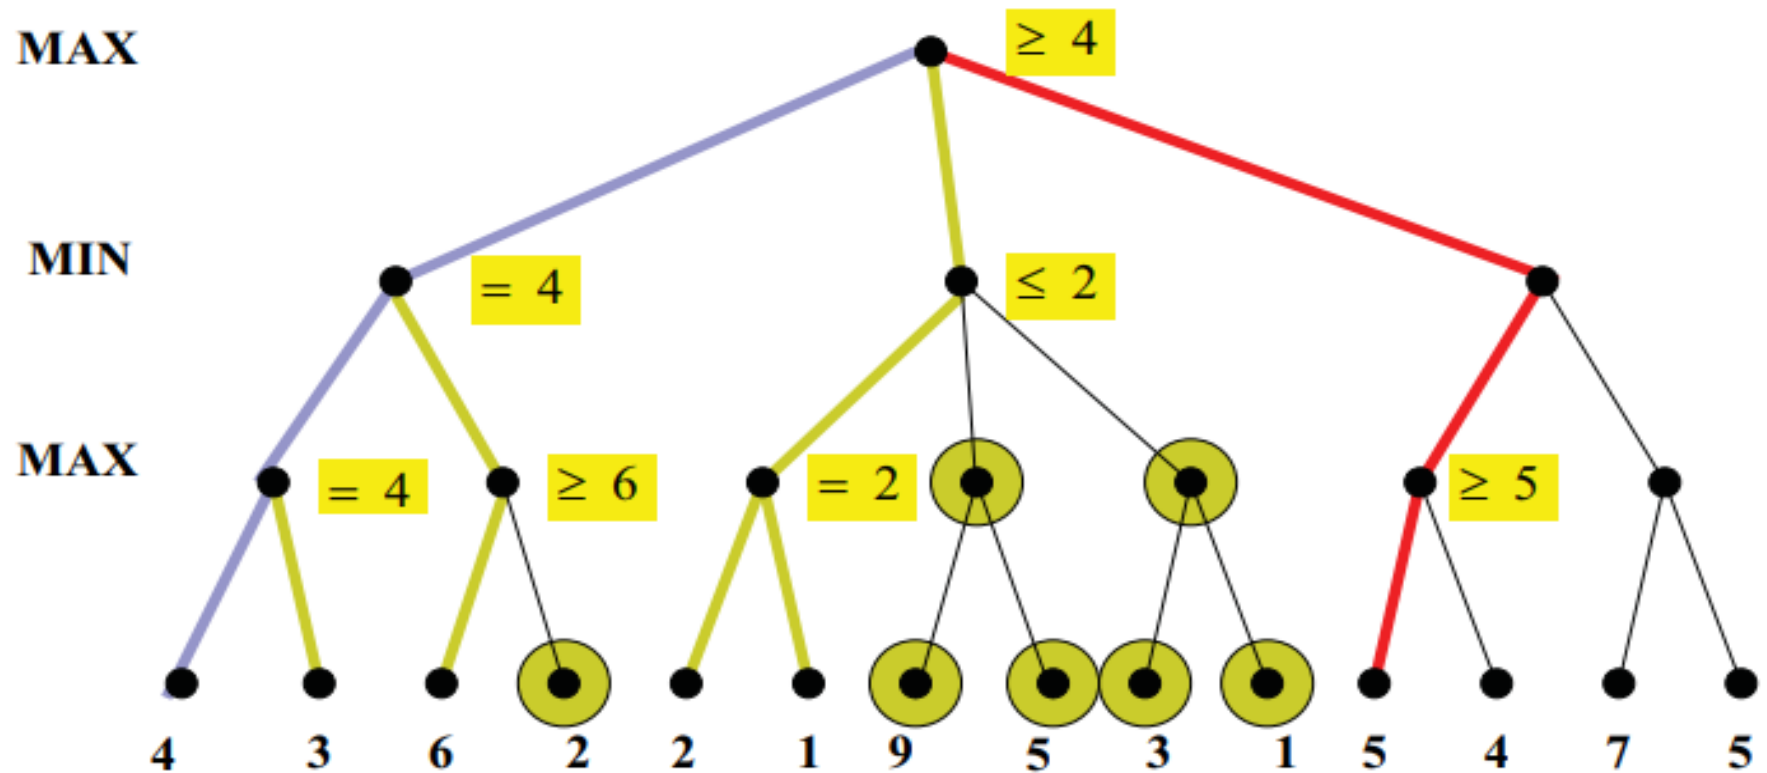
# Alpha beta pruning. Example



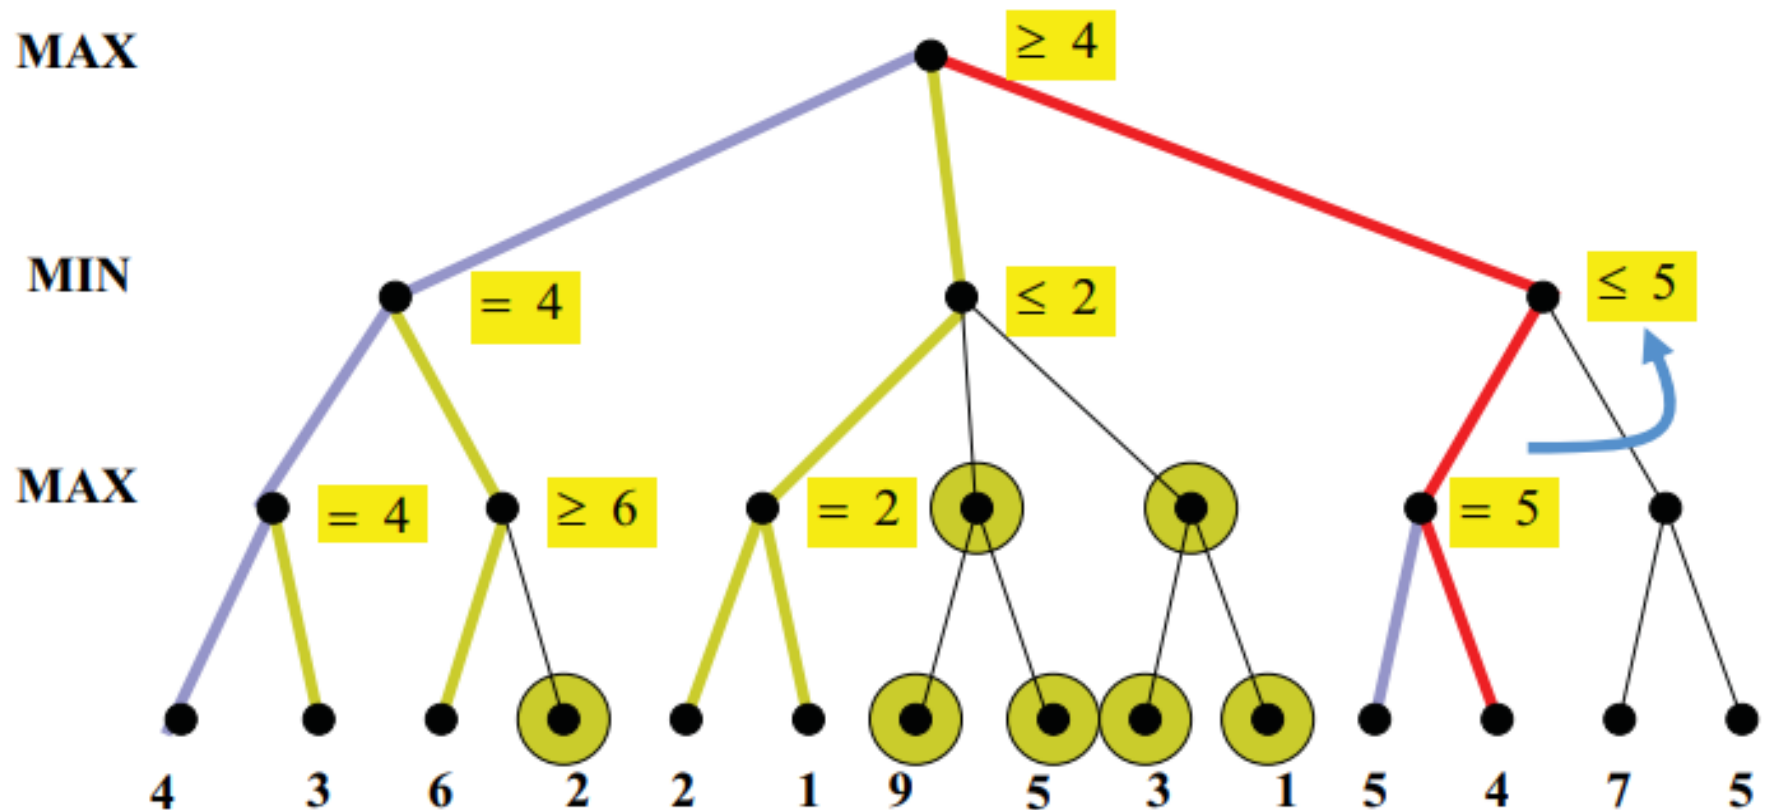
## Alpha beta pruning. Example



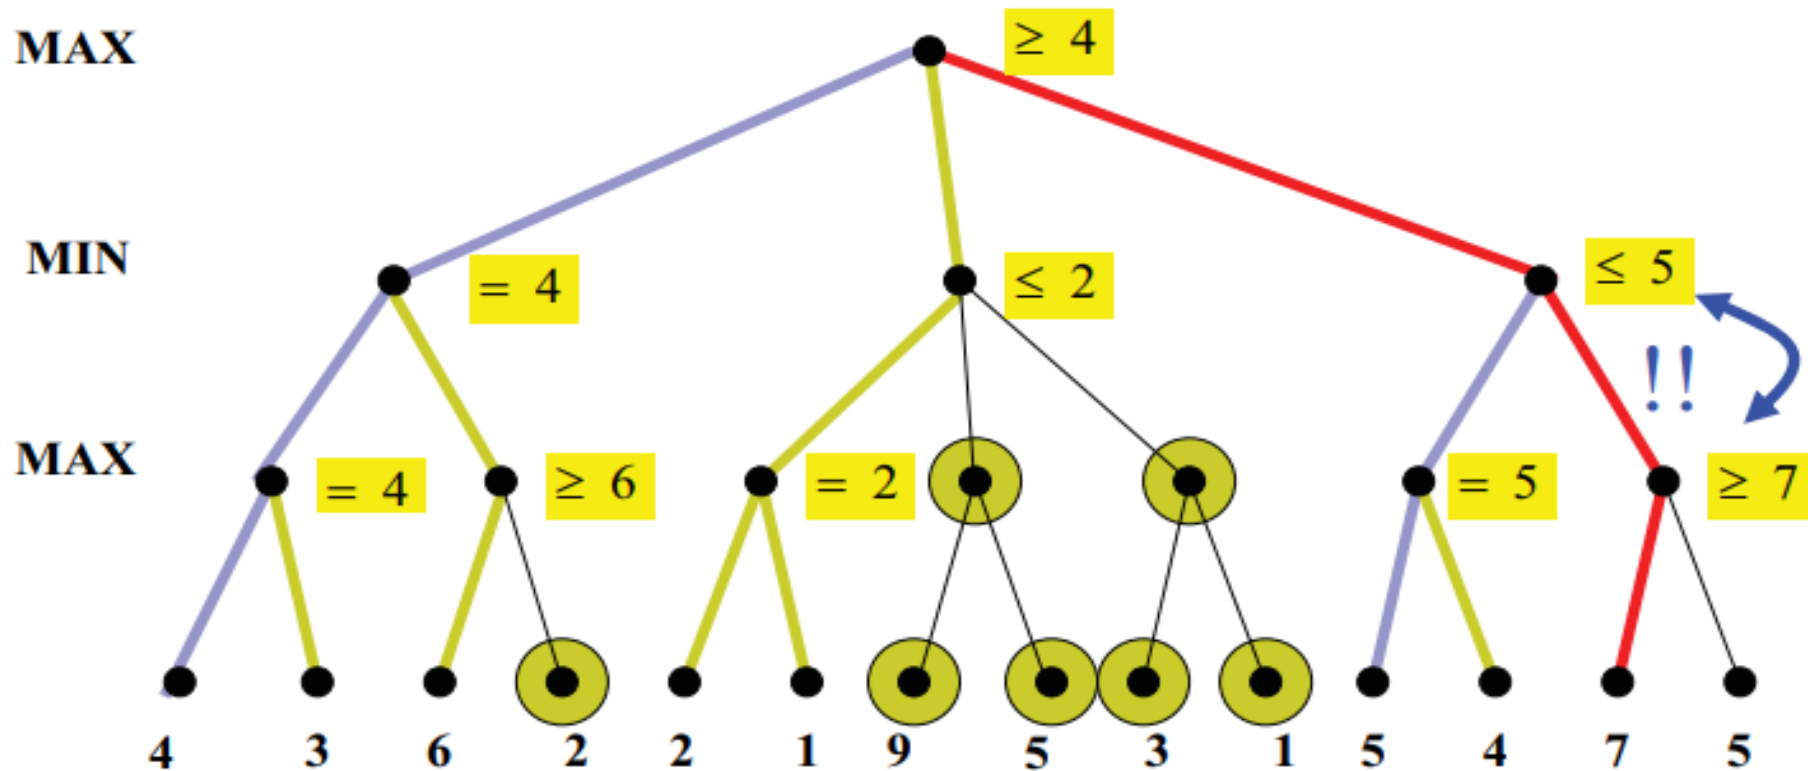
## Alpha beta pruning. Example



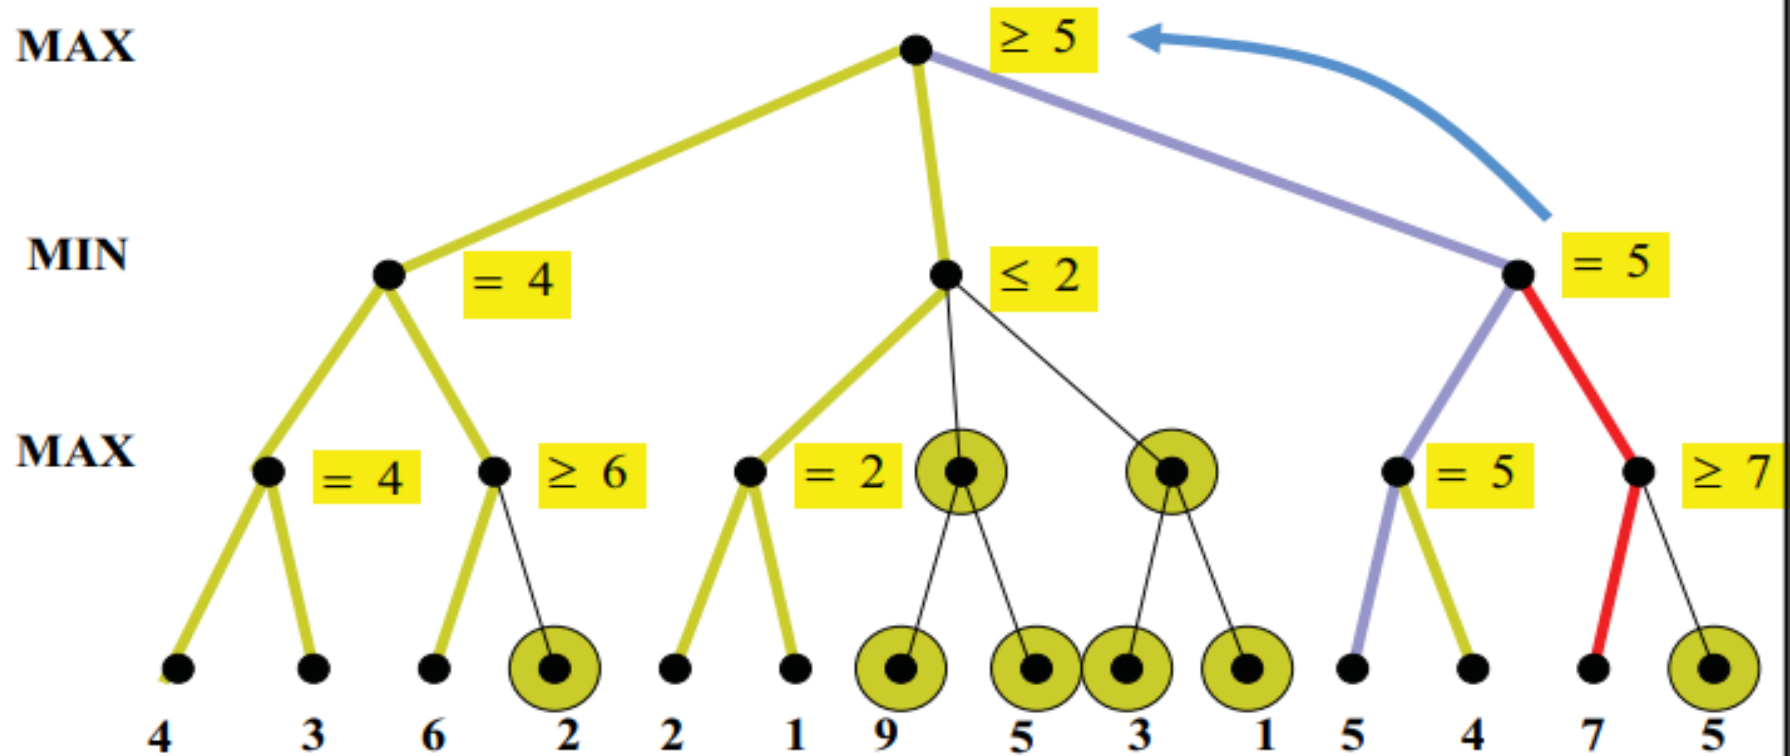
# Alpha beta pruning. Example



# Alpha beta pruning. Example

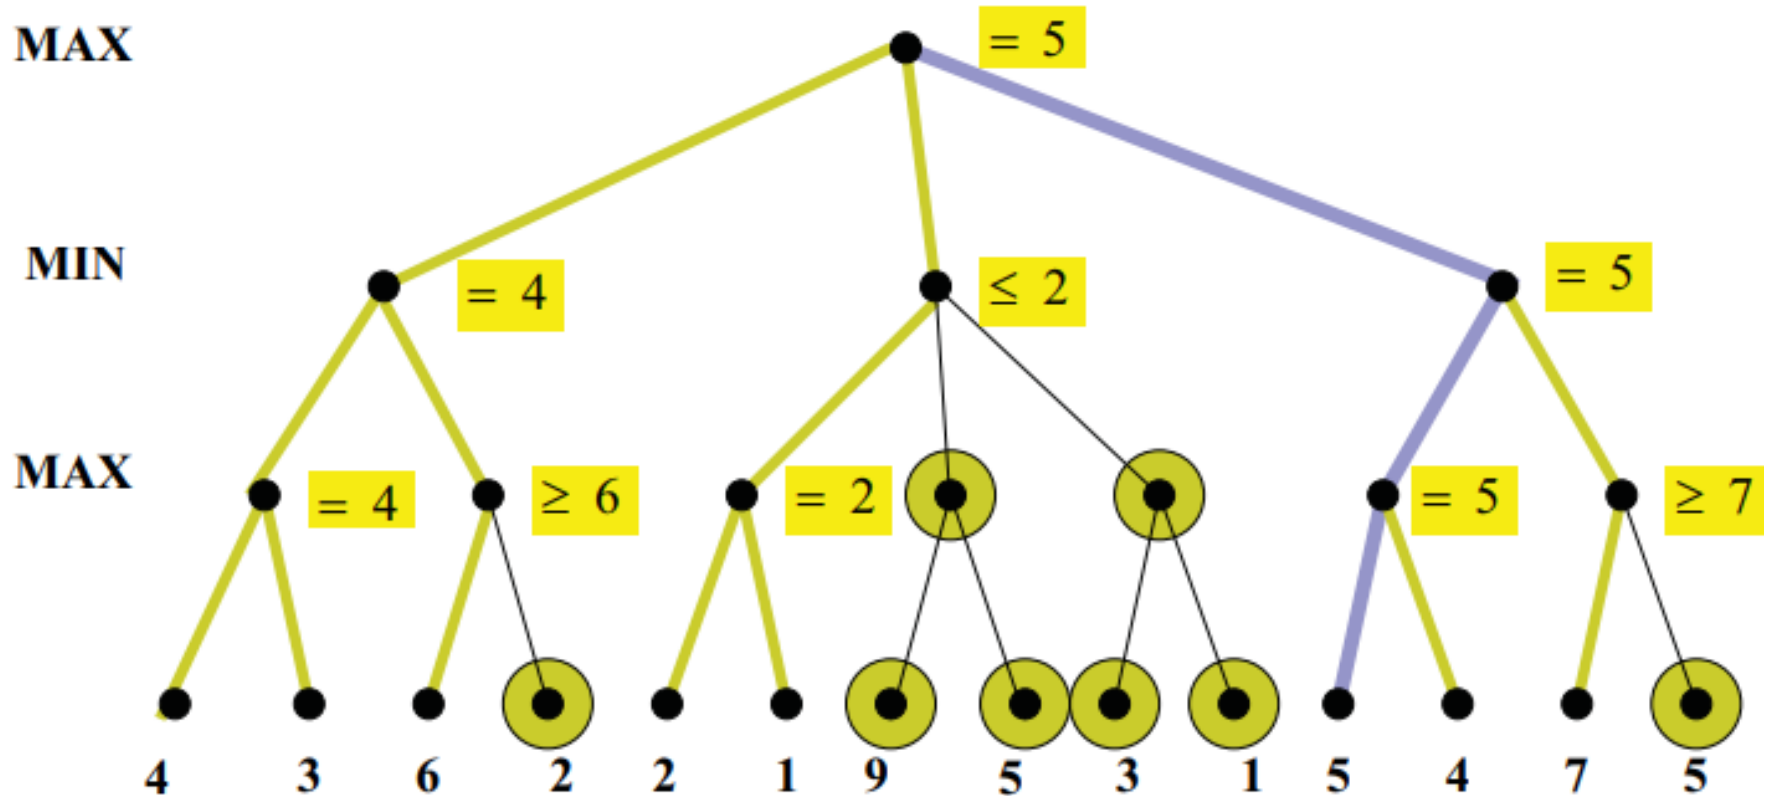


## Alpha beta pruning. Example





## Alpha beta pruning. Example



nodes that were never explored !!!